Difficulty is not fragility: an audited robustness predictor and 88 negative results

Jude Wallis (Independent researcher)

Technical report, AIMO Interpretability Challenge 2026 (NeurIPS 2026, Codabench 16180), 21 October 2026 (UTC); four-page version (draft 31). The full report (draft 29, `report/REPORT_full.pdf` in the repository) is its supplementary material: "Supplement X" points to its section or appendix X. Numbers come from our research ledger (entry in brackets), Figure 1's script or the supplement's audit scripts (Supplement F.7); hidden cases appear as counts only. Number labels are defined in Appendix A.

Abstract

We predict, inside a fixed GPU budget, whether a configuration's accuracy on a math problem survives the organisers' perturbations, and find that difficulty is not fragility. A stronger model's answer agreement reads a problem's difficulty and ranks the 71 labelled public rows at within-configuration AUC 0.929, mostly because 33 of 34 rows on easy templates are robust. On the 37 rows of hard templates the AUC is 0.875, rows that share an agreement but not a label mostly belong to different configurations, and 2.5 times more attempts or a perfect answer key add at most 0.054 AUC and 2 rows. What is left is configuration-specific fragility, and no in-budget attempt to read it passed all its pre-set bars: of 88 negative results, 70 failed and 2 were unclear. Leave-one-template-out our entry gets 60 of 71 public and 18 of 24 fresh rows, a selection null cannot tell its board-chosen layer's gain from coin flips, and a 21-case validation board total varies by about 1.3 cases across server states.

1. Introduction

The AIMO Interpretability Challenge (Štefánik et al., 2026) asks, for each *case*, a (problem, configuration) pair, whether the configuration's accuracy survives perturbation of the problem. The configurations are Qwen3.5-4B, Skywork-OR1-Math-7B, Olmo-3-7B-Think and DeepSeek-R1-0528-Qwen3-8B (the small track), and gpt-oss-120b at low, medium and high effort. A case is robust iff the largest accuracy drop over the organisers' families of perturbed texts is at most 0.10, non-robust iff at least 0.25, and exists only if the configuration solves the original robustly [2]. An entry sees only the original text, runs each case in its own process on one GPU [178, 464] and faces 30 to 50 new problems in the final at about 170 s each [513, 537].

We find that **difficulty is not fragility**: a stronger model's reading of difficulty carries most of the public ranking, but on hard problems one threshold cannot serve every configuration, and no in-budget reading of what remains, configuration-specific fragility, passed all its pre-set bars (Sections 3 and 6). Edited or re-instantiated problems are known to cut accuracy (Mirzadeh et al., 2025; Huang et al., 2025a; Mondorf et al., 2026); whether probes and interpretations predict unseen behaviour is asked by Belinkov (2022), Li et al. (2025) and Huang et al. (2025b). We ask it of a whole entry, with a selection null from adaptive data analysis (Dwork et al., 2015; Blum and Hardt, 2015; Roelofs et al., 2019).

2. Data and the entry

Data. 82 public rows over 28 templates, 71 labelled (52 robust) [3, 278, 382, 618]; *fresh rows* are the 24 labelled rows on 12 AIME-2026 templates [446, 509]; 9 of 12 Skywork rows added on 15 October are labelled (7 robust) [617]. The validation boards hold 21 main cases (10 robust) [99] and 14 small cases (7 robust) [270]. The site invites iterating on them; when we did, each upload was a method variant with a public-row reason, never constant answers per configuration or per-case probes, though early ones tested guesses about the hidden sets' structure (Supplement E.6.2) [266, 282, 290]. Board counts per configuration carry hidden-label information [329b], which we used while choosing three routes (Supplement A.2); Section 4 audits that use.

The entry reads behaviour only. *Stand-in* (Qwen, DeepSeek, gpt-oss low and medium cases): gpt-oss-120b at low effort makes 4 to 32 attempts; their *agreement* (the probability that two attempts that ended on their own agree), a Qwen3.5-4B quick finish and an effort code enter a logistic rule with 10 coefficients (`decision.json`; Supplement F.4) [272b]. *Own runs* (Skywork, Olmo): robust iff at least 1 (Skywork) or 2 (Olmo) of 3 attempts of the judged model end on their own inside a 9,000- or 10,000-token cap [269, 339]. *Prior* (gpt-oss high): robust, as 13 of 15 public rows are [336]. *The entry* is v95 (main) and v96 (small), rebuilt for the final as v111 and v112 [540]; this paper's evaluated entries are their per-case-safe variants v115 and v116 (Codabench 1003729, 1003727; Section 7) [575, 576].

3. Difficulty is not fragility

Two meanings of difficulty. The judged model's solve rate carries almost none of the label: in our fit of the label recipe the chance of a non-robust label is flat, 0.18 to 0.21, for solve rates from 0.5 to 1.0 (computed) [416]; on T13 every small model solves the problem and every one is labelled non-robust [197, 266]. Difficulty read by a stronger model carries far more: the stand-in's agreement at K 32 ranks the 71 public rows at within-configuration AUC 0.929 [611].

That reading saturates on hard templates (Figure 1; computed, post hoc, one server read [440]). Its AUC comes mostly from the 10 easy templates, where 33 of 34 rows are robust. On the 17 *hard* templates, which a 2,000-token reader never finishes (37 rows, 18 non-robust) [267b, 278], it is 0.875: 7 of 56 same-configuration pairs are misordered (four on gpt-oss high, two on low, one on Skywork), 2 of 41 on the stand-in's four configurations; between agreement 0.56 and 0.82, 8 robust and 9 non-robust rows mix, but 69 of their 72 pairs cross configurations: one threshold cannot serve every configuration. 80 pooled attempts give 0.929 and their accuracy against the answer key 0.911, inside the 0.875 to 0.955 of single held reads; leave-one-template-out (LOTO), the key fixes 2 hard rows [611]. Decisions built on agreement get 0.743 of hard rows right (LOTO accuracy, mean of six reads) [278].

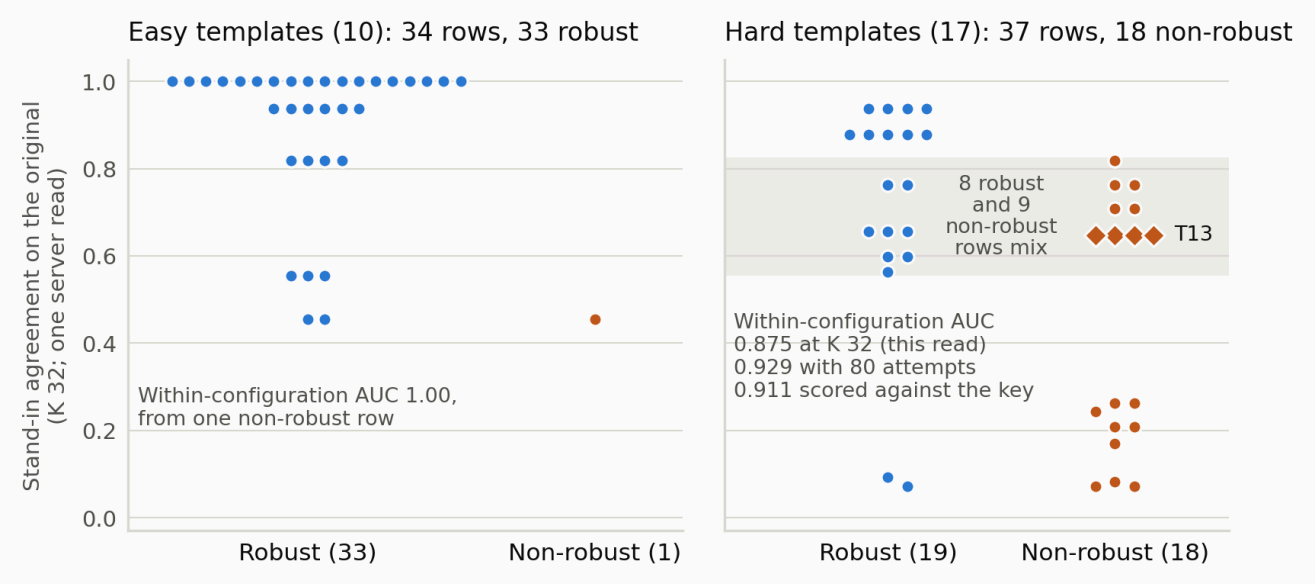


Figure 1. The stand-in's agreement on each labelled public row (one server read, K 32), by label, on easy and hard templates. Diamonds: T13. Shading: the range where both labels occur. Script `figures/make_agreement_by_label_v30.py`.

What fragility can look like. On two harsher expert-style T10 texts, gpt-oss low's wrong answers are exactly those of a mutant of the correct method that skips a triangle bound, equivalent at the original numbers (850 = 850; Papadakis et al., 2015) [446]. gpt-oss medium shares the slip yet is labelled robust, so our texts are harsher than the organisers' [318].

4. A generalisation audit

What was fitted, and when. The only label-fitted part, `decision.json`, was fitted on the 71 public rows on 5 October [113], after board uploads began (2 October) [20] but before the board-tested changes of 5 to 11 October; a LOTO refit changes 1 of 71 answers [509]. Three choices of 10 to 11 October (the gpt-oss high prior and both own runs) decide 8 of 21 main and 6 of 14 small cases; about ten hand-set values (caps, K, thresholds) cannot be cross-validated. Table 1 gives out-of-sample accuracy.

Table 1. The entry's rules on public rows (same server read), LOTO with `decision.json` refitted per held-out template (template-bootstrap 95% intervals; the 80-row line pre-registered [618]), the board this predicts (model-based) and the board observed (OFFICIAL, one draw each) [471, 560, 564].

Rows, LOTO	Accuracy	Predicted board (90% range)	Board
All 71 (in-sample 61) [509]	60/71 = 0.845 [0.75, 0.92]	main 17.0 (13 to 20)	16 (v95), 17 (v111)
24 fresh AIME-2026 rows [509]	18/24 = 0.750 [0.59, 0.91]	main 15.7 (11 to 20)	16, 17
34 small-model rows [509]	0.817 at the board's class mix	small 11.4 (8 to 14)	10 (v96), 9 (v112)
80 rows, 8 reads pooled [618]	68.75/80 = 0.859 (65 to 71 per read; balanced 0.826; always-robust 0.738)	—	—

The board-chosen layer against a selection null. Keeping a change only when the board rises gains +3.9 in expectation (90% range 1 to 8) if every changed answer is a coin flip and +7.3 (3 to 12) if each is right 70% of the time (model-based). The three routes kept +3, inferred from board counts (Supplement A.2) [509], inside both ranges: the test cannot separate them, and **the board-chosen layer adds no demonstrable accuracy**. Moving Skywork's cases to the stand-in gains on the original public rows in four of five reads [543], and three of the six LOTO fresh-row errors sit on board-chosen routes [509]. On the 9 later Skywork rows the own run gets 9 and 7 on two reads, the stand-in about 7.3 over eight [617, 618]: they cannot separate the routes.

A frozen prediction. We expect v111 to get **about 16 of 21 main cases right on new problems** (90% range 11 to 20), not our best draw's 19 [393, 509]. Guess ranges for final accuracy, frozen on 17 October [543] (assumptions, hash and their rule for the evaluated v115 and v116: Appendix A), are 0.72 to 0.82 for v111 and 0.62 to 0.76 for v112; Section 5's noise alone gives them only about 0.4 to 0.5 and 0.55 to 0.7 coverage (computed). **Outcome.** The final had not been scored when this report was submitted; we will report the evaluated entries' accuracy against these ranges, read with the coverage note, whichever way it falls.

5. How much a board score moves

From our own logs, without hidden labels (computed) [585]: a reading repeats exactly within one server state (23 of 23 same-problem groups), but across states a stand-in call flips with probability about 0.26 (90% interval 0.14 to 0.45). One board total thus varies by about 1.3 cases on main and 1.0 on small, matching four earlier draws of the same rules (19, 17, 18, 16) [393]. **A 21-case board cannot order entries whose readings resample and that differ by two cases or fewer**: 17 against 15 arises with probability 0.41 with no real difference. On 30 to 50 new problems, reading noise and problem sampling give about 0.07 to 0.09 accuracy per run.

6. Negative results

Of 88 negative results (Supplement D), **70 failed a bar written before their data** in our dated, author-controlled ledger (no external timestamp), 2 were unclear, 13 had no bar or a post hoc one, and 3 were not built or run; some are engineering levers or board bets. In Table 2, a difficulty variant adds nothing over agreement and no fragility reading passed all its bars.

Table 2. Seven informative readings (server: contest server; local proxy: 4-bit weights on a Mac). Ledger by row: [394], [499, 506, 516, 555], [614], [447, 457, 461, 470, 477, 479, 491], [521, 523, 528], [619], [302, 309].

Reading	Measures	Bar, set first	Result	Verdict
Probe on Qwen3.5-4B prompt states (server)	internal state	balanced accuracy ≥ 0.80	0.669 (stand-in 0.844)	failed
Probes for the rating below: Qwen3.5-4B states on text, solution, thinking (local proxy); gpt-oss-120b (server)	internal state	Spearman ≥ 0.60, 28 templates	0.277, 0.454, 0.304; 0.030	failed
Token confidence for agreement, four forms (Fu et al., 2025; server)	difficulty	+2 rows of 71; +1 on a second read	−3, 0, −12, −9; second read −1, 0, −7, −3	failed
Five blinded Claude raters: would new numbers make a dormant condition bind? Five in-hour versions	fragility	AUC ≥ 0.70, 37 hard rows; one bar per version	0.646; none of five passed	failed
STD: own steps against own brute force on new numbers; three blind writers	fragility	11 of 12 7B rows; AUC ≥ 0.80; Spearman ≥ 0.40	9, 11, 10 of 12; AUC 0.724, 0.900, 0.789	each failed a bar
Premise reliance; instance-binding depth (local proxy)	fragility	Spearman with the rating ≤ −0.30; sabotage loses	+0.150 (11 of 34 problems defined); −0.038 (27); sabotage did not lose	failed
Answer-keeping organiser-style rewrites, gpt-oss low (server, 7 rows)	fragility	drop ≥ 0.20 on non-robust rows	+0.03, +0.10; AUC 0.30	unclear (0.10 is the boundary)

At these sizes 80% power needs a true Spearman of about 0.51 (28 templates) or an AUC of about 0.77 (19 against 18 rows) (Supplement D): "failed" means not shown, not absent. One read of the 80-row bench scores 65 to 71 [618], so +1- and +2-row gates sit within read-to-read spread [616].

Probes held signal inside their 325-problem training pool (nested Spearman 0.346) that did not transfer to the 28 templates [499, 516]: **we make no White-Box claim**. **STD** programs a Claude agent wrote after reading the raters' notes scored 12 of 12 7B rows, blind ones 9 to 11 [482, 489]: **a reading computed through a solver's own steps needs a blind replication.**

7. Efficiency under a 170 s per-case budget

Our decode paths share one K/V read across a KV head's query heads and write K and V in place [360, 186]; in pre-registered server benches a decode run takes 0.42x (Skywork) and 0.48x (Olmo) of the default wall time, and a gpt-oss-120b step 54.9 against 225.2 ms [365, 193]. v115 and v116 answered every board case in at most 148 s, under the final's 170 s [575, 576].

8. Limitations and lessons

Limitations. Reading noise is measured coarsely (3 server states) [585]; fresh rows were seen while hand-set values were chosen; the own runs rest on thin, post hoc support (3 Skywork and 9 Olmo rows when chosen) [339, 387]; a Claude agent wrote the rating rubric after seeing the labels [447, 480]; raters and blind writers are Claude subagents; the final is mocked, not run [538]. **Lessons.** Split a reading's score by where the label varies; refit label-fitted parts LOTO and compare board gains with a selection null (`audit.py`); set bars above read-to-read noise.

AI Use Statement

Generative AI was used for most of this work. We used Anthropic's Claude through Claude Code, a tool in which Claude writes and runs code, reads and edits files and starts further Claude instances (subagents). The Claude models used were Claude Opus 5.5, Claude Opus 5 and Claude Fable 5.1. Claude was used for:

- **ideation**: research directions, hypotheses and the design of each pre-registered test (bars, controls, outcome branches);
- **coding**: the entries, decode paths, mocks, fault injection, the analysis and audit scripts, the release folders and the figure scripts; Claude also created and updated the public GitHub repository under the author's account, with the author's permission [632, 634];
- **research data**: Claude agents wrote the perturbed problem texts (1,976 texts) [302], the LATENT rating rubric [447, 480] and the solution programs of Section 6 [482];
- **experiments and analysis**: local runs, the statistics and the research ledger; Claude also operated the competition site in the author's logged-in browser session, to read server logs and, from 5 October, to upload entries and research runs, among them v95 and v96 [469] and the evaluated v115 and v116 [574], under permissions the author gave in chat (at first for one night [125], later for a session or a named file [239, 503], and standing from 18 October [556, 601]). One upload went up under a standing go-ahead before the author had answered a question about it [288], and one file was uploaded twice by mistake [295];
- **literature**: search and reading (statuses in the References);
- **writing**: this report, its supplement and its consistency checks; Claude instances also reviewed drafts and code.

Claude instances were also the five blinded raters and the three blind program writers of Section 6 (a limitation; Section 8). No Claude model runs inside the entries, which use only the open models of Section 1.

Tasks done without generative AI. The author set the goals and limits of the work; gave the permissions for uploads and compute in chat, mostly as standing permissions; provided the competition, GitHub and other accounts; sent this project's messages to the organisers (Claude drafted some and sent none); made the final decisions, including which entries are evaluated with this paper (Claude recommended them); and submits this report.

Review and responsibility. The author reviewed all AI-assisted work and takes full responsibility for the final content of this paper, including AI-assisted text, claims, code and other artefacts.

References

Status: *read* (we read the paper), *abstract* (abstract page only), *listed* (details from the organisers' proposal's bibliography; we did not open the work), *record* (authors and venue checked in the publisher's record; abstract not seen).

- Belinkov (2022). Probing classifiers: promises, shortcomings, and advances. Computational Linguistics 48(1). *Listed*.
- Blum and Hardt (2015). The Ladder: a reliable leaderboard for machine learning competitions. ICML 2015. *Abstract*.
- Dwork et al. (2015). The reusable holdout: preserving validity in adaptive data analysis. Science. *Abstract*.
- Fu, Wang, Tian and Zhao (2025). Deep think with confidence (DeepConf). arXiv 2508.15260. *Abstract* [583].
- Huang et al. (2025a). MATH-Perturb: benchmarking LLMs' math reasoning abilities against hard perturbations. ICML 2025; arXiv 2502.06453. *Abstract*.
- Huang, Tao, Icard, Yang and Potts (2025b). Internal causal mechanisms robustly predict language model out-of-distribution behaviors. ICML 2025. *Listed*.
- Li, Kaufmann, Wattenberg, Alvarez-Melis and Saphra (2025). Can interpretation predict behavior on unseen data? arXiv 2507.06445. *Listed*.
- Mirzadeh et al. (2025). GSM-Symbolic: understanding the limitations of mathematical reasoning in large language models. ICLR 2025. *Listed*.
- Mondorf, Bell, Dodge and Hupkes (2026). LPDS: evaluating LLM robustness through logic-preserving difficulty scaling. arXiv 2605.15393. *Read*.
- Papadakis et al. (2015). Trivial compiler equivalence: a large scale empirical study of a simple, fast and effective equivalent mutant detection technique. ICSE 2015. *Record*.
- Roelofs et al. (2019). A meta-analysis of overfitting in machine learning. NeurIPS 2019. *Abstract*.

- Štefánik, Mondorf, Waldis et al. (2026). AIMO Interpretability Challenge (the organisers' proposal). arXiv 2607.13899. *Read*.

Appendix A. Supplementary material and reproducibility details

Reproducibility. v111, v112, v115, v116, v121 and v122 rebuild byte for byte from the repository's root folder (v115 / v116 and v121 / v122 from its folders `v115_v116/` and `v121_v122/`; md5s below; every command in its README) [549, 589, 627]. **Training data**: `training/build_calib.py` rebuilds `calib.json` from the organisers' public parquet (Hugging Face `aimo-interp/train-main-v2`, revision `23791d9b`), and `training/refit_low16.py` refits `decision.json` from it and our four public-template server logs (`training/runs/`). **Board scores**: each zip runs unchanged in the organisers' per-case harness; a rerun in another server state varies by about 1.3 cases on main and 1.0 on small (Section 5) [585]. **Licence**: code under the MIT License; report text under CC BY 4.0. **Repository**: https://github.com/EconLearn/aimo-interp-difficulty-vs-fragility.

Evaluated entries. The entries evaluated with this paper are Codabench 1003729 (v115, main) and 1003727 (v116, small) [574]; by the rule fixed before any final result [568], the frozen ranges apply to them if they alone are scored. The other rows are listed for reproducibility (md5 of each zip; board OFFICIAL, one draw each, with its Codabench submission id):

Role	File	md5	Board
The entry rebuilt for the final, main; small	`aimo_v111_main.zip`; `aimo_v112_small.zip`	9686c9cec06ec4dc40f9b728f7ef17db; 8b404890d71bc5483ccce18f88887636 [540]	17 of 21 (1000991); 9 of 14 (1000212) [560, 564]
Per-case-safe variants (own runs stop at 160 s)	`aimo_v115_main.zip`; `aimo_v116_small.zip`	ad191e2e2262b0a918ec1aaf83bcf9c7; d3ada3a4b8ff05b5876582246008bbf9 [573]	15 of 21 (1003729); 9 of 14 (1003727) [575, 576]
The same with every reading's seed shifted	`aimo_v121_main.zip`; `aimo_v122_small.zip`	22ced3169a4d73b5cfe9b7ceb0d815b6; 4de6852eee8d4c3e6ccb07576c422a31 [588]	15 of 21 (1004377); 9 of 14 (1004381) [608]

Frozen prediction (Section 4). Assumed: a roughly balanced class mix, a validation-like configuration mix and a server at most 1.2x slower than our mock, with 170 s as an average (a per-case limit leaves v111 15 of 21 answered [573]). By a rule fixed before any final result [568], the forecast holds on a track if the final accuracy of v111 (main) or v112 (small), or of another nominee if only that one is scored, lies in its range, so the ranges apply to v115 and v116 if they alone are scored. The prediction's ledger row (`work/miners18/FROZEN_PREDICTION_ledger543.txt`) has sha256 `ffd7c3fcc370631dbc7501bd713444df47f067ede871a672c15ead8aa91a2ec6`, computed on 18 October after the board runs of that day, with no external timestamp [566].

Supplementary material. The full report (draft 29, 38 pages; `report/REPORT_full.pdf` and `report/REPORT_full.md` in the repository) holds every detail this version omits; its supplement folder `report_supplement/` holds the full text of its longer appendix parts. Where to look: the audit in full, Supplement E.13, and the board draws against it, the supplement's Figure 1; every board upload, Supplement A; board-feedback use, Supplement A.2; the 88 negative results with their bars, Supplement D and E.6; the worked T10 example, Supplement E.12; decode paths, Supplement E.3; the final-shape mocks and the pre-mortem, Supplement E.15 and E.17; token confidence, the 80-row rescore, the key-headroom check and the fragility vectors, Supplement E.18; the decision rule's coefficients, Supplement F.4; scripts, including `audit.py`, Supplement F.7. Figure 1 is made by `figures/make_agreement_by_label_v30.py` from held logs (lane G's loader and the server read of 15 October); it writes its numbers to `figures/make_agreement_by_label_v30.out`.

Number labels. *Board*: the OFFICIAL hidden validation score, counts only, one draw each. *Public rows*: the labelled public rows (in-sample where our rule was fitted on them). *Server*: a research run on the contest server. *Computed*: offline, from logs or outputs we hold. *Mock*: a rehearsal on a virtual clock (time, coverage and validity, not accuracy). *Model-based*: a simulation. *Local proxy*: a run on a Mac with 4-bit weights standing in for a judged model. *Guess*: a judgement, given as a range.

Ledger entries. Every bracketed number is an entry of our dated research ledger; the full report cites the same entries, except 589 and 627 (release rebuild checks). The ledger holds internal notes and is not published as it stands; a redacted extract of the cited entries is available on request.